



UPPSALA  
UNIVERSITET



# Data Wrangling - Importing and preparing data for analyses

## R for Psychology Research

Marcus Lindskog, Docent  
2019-09-23

# Overview

1. Keep data in a `tibble`
2. Import data with `readr`.
3. Data transformations with `dplyr`
4. Data tranformation with `tidyr`

# What is data wrangling?

- Data wrangling is the process of getting your raw data into a format that can be used for analyses.
- Today we will work with a collection of packages from the `tidyverse` to do that.
- It is important to acknowledge that this is not the only way to do things.
- However, using the `tidyverse` gives you an integrated framework that helps you solve these tasks quite easy.

Keep your data in a tibble

# What is a tibble?

- The `tibble` package introduces a new data structure into R, the **tibble**.
- A tibble is a modern take on the data frame.
- It makes it easier and more consistent to work with tidy data in R.

# Creating a tibble

- A regular `data.frame` can be coerced to tibble with:

```
library(psych)  
library(tibble)  
  
as_tibble(bfi)
```

# Create a tibble from individual vectors

- A tibble can be created from individual vectors with:

```
tibble(  
  x = 1:5,  
  y = 1,  
  z = x^2 + y+1,  
  w = letters[1:5]  
)
```

# Subsetting a tibble

- A tibble can be subsetting by name or position. But it works a bit different than for `data.frame`

```
my_data <- tibble(  
  x = 1:20,  
  y = 1,  
  z = x^2 + y+1,  
  w = letters[1:20]  
)
```

```
# Extract by name (1)
```

```
my_data$w
```

```
# Extract by name (2)
```

```
my_data[["w"]]
```



# Subsetting a tibble

```
# Extract by position
```

```
my_data[[4]]
```

Import data with readr

# readr.

- readr is a packages that turns flat files with rectangular data into data frames (tibble).
- The same functionality can be found in base R, but readr functions generally:
  - Are much faster.
  - Produce tibbles and don't make to many assumptions about your data (e.g., turns strings into factors).
  - Are more reproducible.

# Functions in readr

- `read_csv` - reads comma-delimited files
- `read_csv2()`: reads semicolon-separated files
- `read_tsv()`: reads tab-delimited files
- `read_delim()`: reads in files with any (user specified) delimiter.

# Functions in readr

- **read\_csv - reads comma-delimited files**
- `read_csv2()`: reads semicolon-separated file
- `read_tsv()`: reads tab-delimited files
- `read_delim()`: reads in files with any (user specified) delimiter.

# A note on the working directory

- When reading a file you need to know where it is (i.e., provide a path)
- But it is hard to give a path if you don't know where you are.
- You are in you're **working directory**, which can be found with:

```
getwd()
```

```
## [1] "/Volumes/ipb-users-1/marli361/@Undervisning/@Kurser/FoU - R for Psychology
```

- You can change your working directory with:

```
setwd("path")
```

- Or you can click **More** under the **Files** tab in RStudio.

# Reading that .csv-file!

```
# read_csv("path")  
heights_data <- read_csv("heights_data.csv")  
bfi_data <- read_csv("heights_data.csv")
```

# A bit more specific parsing

- `read_csv()` assumes that the first row contains column names. It also reads all lines, if nothing else is specified. We can change that behavior.

*#To skip one or more lines*

```
bfi_data <- read_csv("bfi_data.csv",  
                     skip = 3)
```

*# To skip lines that begin with a specific character (i.e., comments)*

```
bfi_data <- read_csv("bfi_data.csv",  
                     comment = "#")
```

*#If the data don't have column names in the first row*

```
bfi_data <- read_csv("bfi_data.csv",  
                     col_names = FALSE)
```

*#To provide your own colnames if they are missing*

```
height_data <- read_csv("height_data.csv",  
                        col_names = c("A", "B"))
```



# It is guess work.

- When you read a file with `readr` it tries to guess what data types are in your columns.
- This is good because it makes the function fast, but it can sometimes be problematic.
- If you know what you have in your columns, you can specify that directly in the `read_`-functions.

*#To skip one or more lines*

```
bfi_data <- read_csv("heights_data.csv",  
                     col_types = col(  
                       height = col_integer(),  
                       cubit = col_character()  
                     )
```

# Other file formats.

- There are of course a lot of different file formats you might want to get into R.
- We can't cover them all. However, have a look at the following packages to solve some of your importing needs.
- haven: reads SPSS, STATA, and SAS files.
- readxl: reads Excel files (both .xls and .xlsx).
- DBI along with a database-specific backend (RMySQL, RSQLite, RPostgreSQL) to run queries against database.

# Data transformation with dplyr

# Functions from dplyr:

- Select, filter and arrange your data:
  - `select()`: Select columns from your dataset
  - `filter()`: Filter out certain rows that meet your criteria(s)
  - `arrange()`: Arrange your column data
- Create new variables:
  - `mutate()`: Create new columns by preserving the existing variables
- Summarize that data:
  - `group_by()`: Group different observations together.
  - `summarise()`: Summarise any of the above functions
- Join data with other data frames
  - `join()`: *Perform left, right, full, and inner joins in R*

# filter()

- The `filter()` function subsets a data frame based on a series of criterion

```
heights_data <- read_csv("heights_data.csv")  
#filter(data_frame, expression_to_filter_1, expression_to_filter_2,...)  
filter(heights_data, height == 71)  
filter(heights_data, height == 68, height == 71)  
filter(heights_data, !(height == 68 | height == 71))  
filter(heights_data, height %in% c(68, 70, 71))
```

- `filter()` returns a tibble and the input is left unchanged.

# Piping that filter

- If you are running a large number of manipulations on the same data frame it is clunky to save each intermediate step in a new variable.
- To help you overcome this problem we have the pipe operator `%>%` from the ``magrittr``

```
heights_data <- read_csv("heights_data.csv")  
  
#filter(data_frame, expression_to_filter_1, expression_to_filter_2,...)  
  
filtered_heights_data <- heights_data %>%  
  filter(height == 71)
```

# arrange()

- The arrange() function changes the order of the rows in a data frame

```
bfi_data <- read_csv("bfi_data.csv")  
  
#arrange(data_frame, column_to_arrange_1, column_to_arrange_2,...)  
  
arrange(bfi_data, A1, A2)  
  
#Use descending order instead  
  
arrange(bfi_data, desc(A1), A2)
```

- Of course, filter and arrange can be combined

```
bfi_data <- read_csv("bfi_data.csv")  
  
new_filtered_arranged_data <- bfi_data %>%  
  filter(A1 %in% c(1,2,3)) %>%  
  arrange(desc(A1), A2)
```

# select()

- `select()` helps select a subset of columns from a data frame.

```
bfi_data <- read_csv("bfi_data.csv")  
  
#select(data_frame, column_1, column_2,...)  
  
#select by name  
select(bfi_data, A1, A2, C5)  
  
#select an interval  
select(bfi_data, A1:C5)  
  
#select all but specified columns  
select(bfi_data, -A1, -A2, -C5)  
select(bfi_data, -(A1:C5))
```



# select()'s little helpers

- starts\_with("arn") matches columns beginning with "abc".
- ends\_with("klm") matches columns ending with "klm".
- contains("una") matches names containing "una".

```
bfi_data <- read_csv("bfi_data.csv")  
  
select(bfi_data, starts_with("A"))  
  
#select an interval  
select(bfi_data, ends_with("4"))
```

# Rename a variable

- `rename()` is a useful tool to rename columns.

```
bfi_data <- read_csv("bfi_data.csv")  
rename(bfi_data, A_1 = A1)
```

# Let's combine

```
bfi_data <- read_csv("bfi_data.csv")  
  
new_changed_data_frame <- bfi_data %>%  
  select(starts_with("A")) %>%  
  rename(A_1 = A1, A_4 = A4) %>%  
  filter(A_4 %in% c(3,4,5)) %>%  
  arrange(A_1, desc(A3))
```

# mutate()

- It is **very** often the case that we need to create (add) new variables that are some combination of existing variables or add some new information.
- This is can be done smoothly with `mutate()` and `transmute()`

```
bfi_data <- read_csv("bfi_data.csv")

#mutate(data_frame, new_var_1, new_var_2,...)
# mutate adds a new variable, and keeps all the old ones.
mutate(bfi_data,
  A = A1+A2+A3+A4,
  C = (C1*C2)/(C3*C4)
)

# transmute only keeps the new variables
transmute(bfi_data,
  A = A1+A2+A3+A4,
  C = (C1*C2)/(C3*C4),
  D = as_factor(rep(c("A", "B", "C", "D"), each = n()/4))
)
```

# mutate()'s little helpers.

- Read p. 56-58 for some functions that can be very helpful when using mutate and transmute.
- Other useful versions of mutate is mutate\_all, mutate\_if, and mutate\_at. Google them to find out exactly how they work.

# Let's combine

```
bfi_data <- read_csv("bfi_data.csv")

new_changed_data_frame <- bfi_data %>%
  select(starts_with("A")) %>%
  rename(A_1 = A1, A_4 = A4) %>%
  filter(A_4 %in% c(3,4,5)) %>%
  mutate(K = (A_1 + A_4)/A3) %>%
  arrange(K)
```

# summarize() and group\_by()

- We often want to create summaries of our data.
- summarize() collapses data into a single row.

```
bfi_data <- read_csv("bfi_data.csv")  
  
#summarize(data_frame, summary)  
summarize(bfi_data,  
          a1_mean = mean(A1, na.rm = TRUE),  
          count = n(),  
          c1_na = sum(is.na(C1))  
          )
```

# Grouped summaries

- To get grouped summaries, first use `group_by()`
- `summarize()` collapses data into a single row.

```
bfi_data <- read_csv("bfi_data.csv")

group_by(bfi_data, A1) %>%
  summarize(a2_mean = mean(A2, na.rm = TRUE),
            count = n(),
            c1_na = sum(is.na(C1))
            )
```



# Let's combine

```
bfi_data <- read_csv("bfi_data.csv")

new_changed_data_frame <- bfi_data %>%
  select(starts_with("A")) %>%
  rename(A_1 = A1, A_4 = A4) %>%
  filter(A_4 %in% c(3,4,5)) %>%
  mutate(K = (A_1 + A_4)/A3) %>%
  arrange(K) %>%
  group_by(A2, A3) %>%
  summarize(mean_k = mean(K, na.rm = TRUE),
            count = n(),
            se_age = sd(age, na.rm = TRUE)/sqrt(sum(!is.na(K))))
```

# Data transformation with `tidyr`

# tidyr

- tidyr has the functions we need to rearrange data into different formats.
- It was originally designed to get data into a **tidy** format, but you can of course use it to get your data into any format you need.

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20095360
Brazil	1999	31737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	128042583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20095360
Brazil	1999	31737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	128042583

observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20095360
Brazil	1999	31737	172006362
Brazil	2000	80488	174604898
China	1999	212258	1272915272
China	2000	213766	128042583

values

# Important functions in `tidyr`

- `gather()`: Gathers multiple columns and converts them into key-value pairs.
- `spread()`: Takes two columns and spreads them into multiple columns.
- `separate()`: Helps in separating or splitting a single column into numerous columns
- `unite()`: Works opposite to the `separate()` function. Combines two or more columns into one

# Gathering

```
head(table4a,3)
```

```
## # A tibble: 3 x 3
##   country    `1999` `2000`
##   <chr>      <int> <int>
## 1 Afghanistan    745   2666
## 2 Brazil        37737  80488
## 3 China         212258 213766
```

- Here we have a data frame with *values* (1999, 2000) in variable names.

# gather()

- We can collect them into a variable by using gather()

```
table4a %>%  
  gather(2:3, key = "year", value = "population")  
  
table4a %>%  
  gather(c("1999", "2000"), key = "year", value = "population")  
  
table4a %>%  
  gather("1999", "2000", key = "year", value = "population")
```

# Spreading

- `spread()` does the opposite of `gather()`
- It takes observations that are scattered into multiple rows, and puts them in columns
- What is the problem with this data frame?

```
head(table2, 6)
```

```
## # A tibble: 6 x 4
##   country      year type      count
##   <chr>      <int> <chr>    <int>
## 1 Afghanistan  1999 cases      745
## 2 Afghanistan  1999 population 19987071
## 3 Afghanistan  2000 cases      2666
## 4 Afghanistan  2000 population 20595360
## 5 Brazil       1999 cases      37737
## 6 Brazil       1999 population 172006362
```

```
table2 %>%
  spread(key = type, value = count)
```



# separate()

- `separate()` pulls apart one column into multiple columns.

```
table3
```

```
## # A tibble: 6 x 3
##   country      year rate
## * <chr>      <int> <chr>
## 1 Afghanistan  1999 745/19987071
## 2 Afghanistan  2000 2666/20595360
## 3 Brazil       1999 37737/172006362
## 4 Brazil       2000 80488/174504898
## 5 China        1999 212258/1272915272
## 6 China        2000 213766/1280428583
```

```
table3 %>%  
  separate(rate, c("cases", "population"), sep = "/")
```

```
## # A tibble: 6 x 4  
##   country      year cases  population  
##   <chr>      <int> <chr>    <chr>  
## 1 Afghanistan 1999  745    19987071  
## 2 Afghanistan 2000 2666    20595360  
## 3 Brazil      1999 37737   172006362  
## 4 Brazil      2000 80488   174504898  
## 5 China       1999 212258  1272915272  
## 6 China       2000 213766  1280428583
```

# unite()

- unite() is the inverse of separate()

```
table3 %>%  
  separate(rate, c("cases", "population"), sep = "/") %>%  
  unite(rate, cases, population, sep = "-")
```

```
## # A tibble: 6 x 3  
##   country      year rate  
##   <chr>      <int> <chr>  
## 1 Afghanistan  1999 745-19987071  
## 2 Afghanistan  2000 2666-20595360  
## 3 Brazil       1999 37737-172006362  
## 4 Brazil       2000 80488-174504898  
## 5 China        1999 212258-1272915272  
## 6 China        2000 213766-1280428583
```

That's all folks!